

// GITOPS REPO STRUCTURES AND PATTERNS

Johannes Schnatterer, Cloudogu GmbH

 @schnatterer@floss.social

 in/jschnatterer

 @jschnatterer

Version: 202311130946-0f3e43b

Categories of patterns

- **Operator deployment**: GitOps operators ↔ Clusters/Namespaces
- **Repository**: How many repos?
- **Promotion**: How to model environments/stages?
- **Wiring**: Bootstrapping operator, linking repos and folders

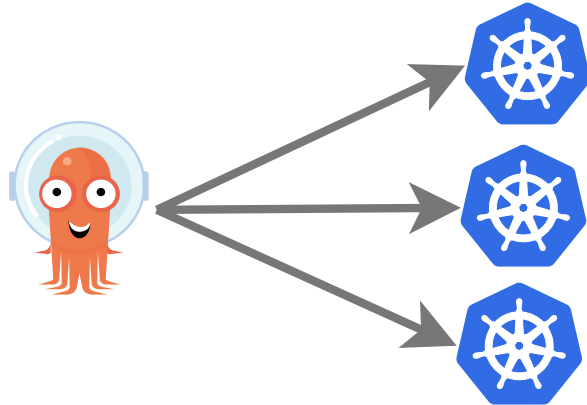
GitOps Operator deployment patterns

How many GitOps operators per cluster?

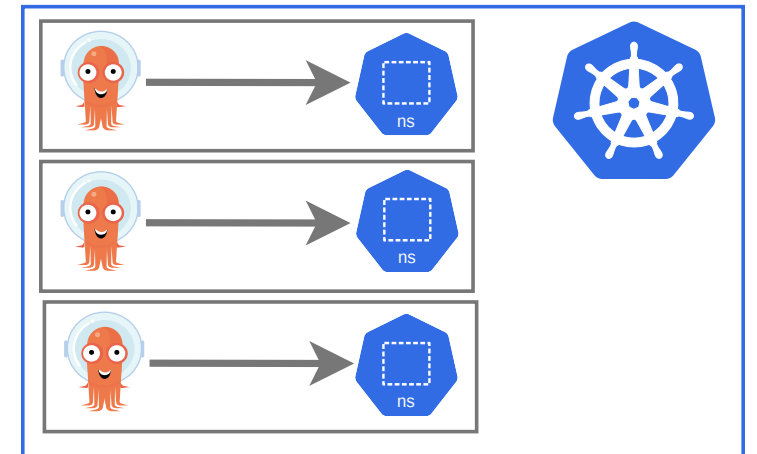
Instance per Cluster



Hub and Spoke



Instance per Namespace



Repository patterns

How many config repos?

- **Monorepo** (opposite: polyrepo)
- **Repo per Team** / Tenant
- **Repo per App**
 - Repo Separation
 - Config replication
 - Repo pointer
 - Config Split
- **Repo per environment** 🕒

💡 Can be mixed 🖨️

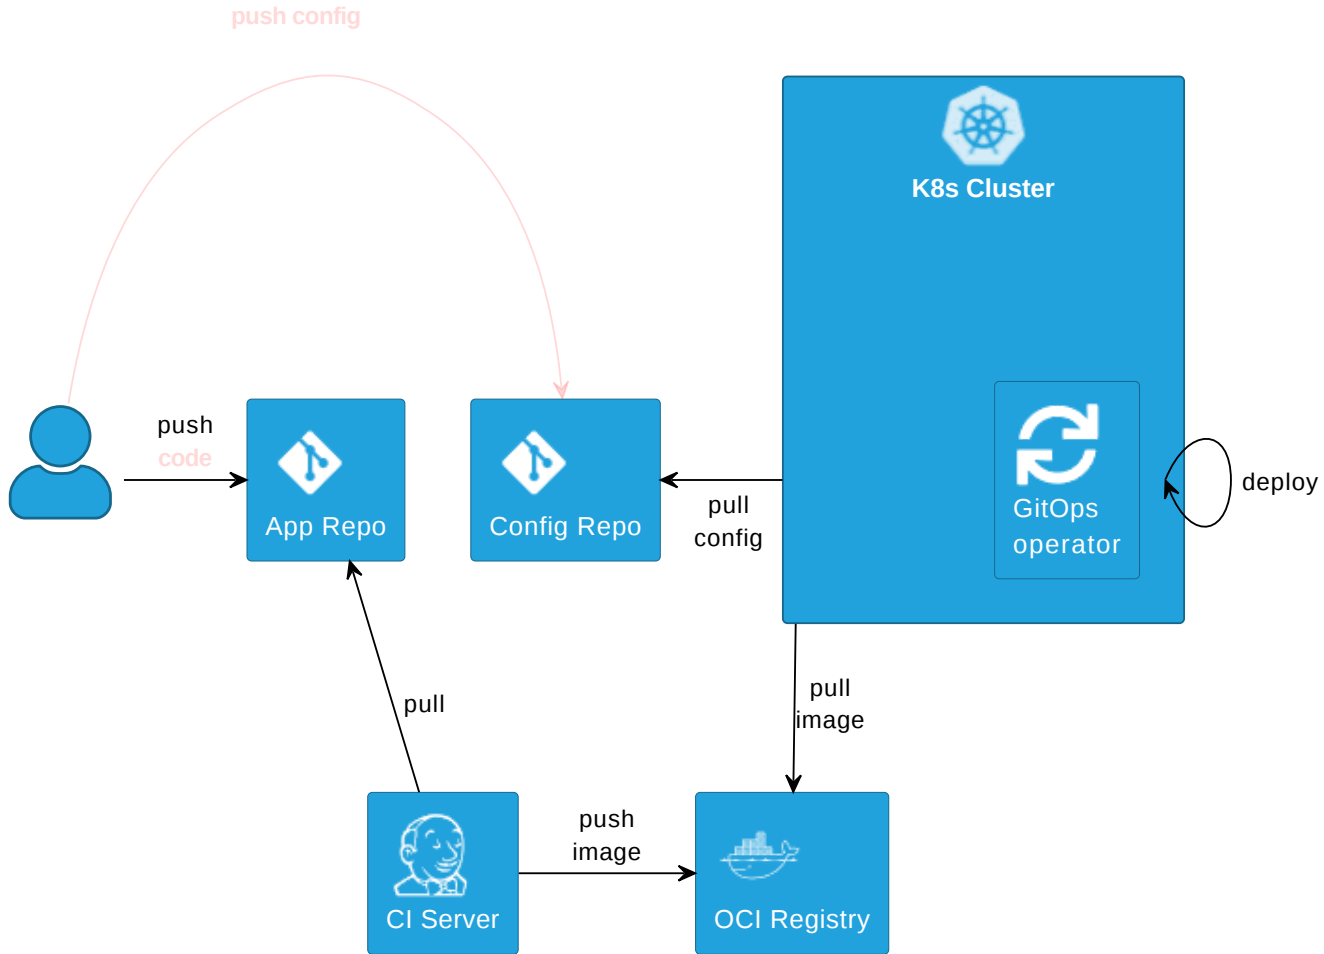
Repository types

	Config repo	App repo
Content	Config/Manifests/YAMLs (IaC)	Application source code
Synonyms	<ul style="list-style-type: none">• GitOps repo• Infra repo• Environment repo• Payload repo	<ul style="list-style-type: none">• Source code repo• Source repo

Example



Repo Separation



Recommendation: Keep config separate from code

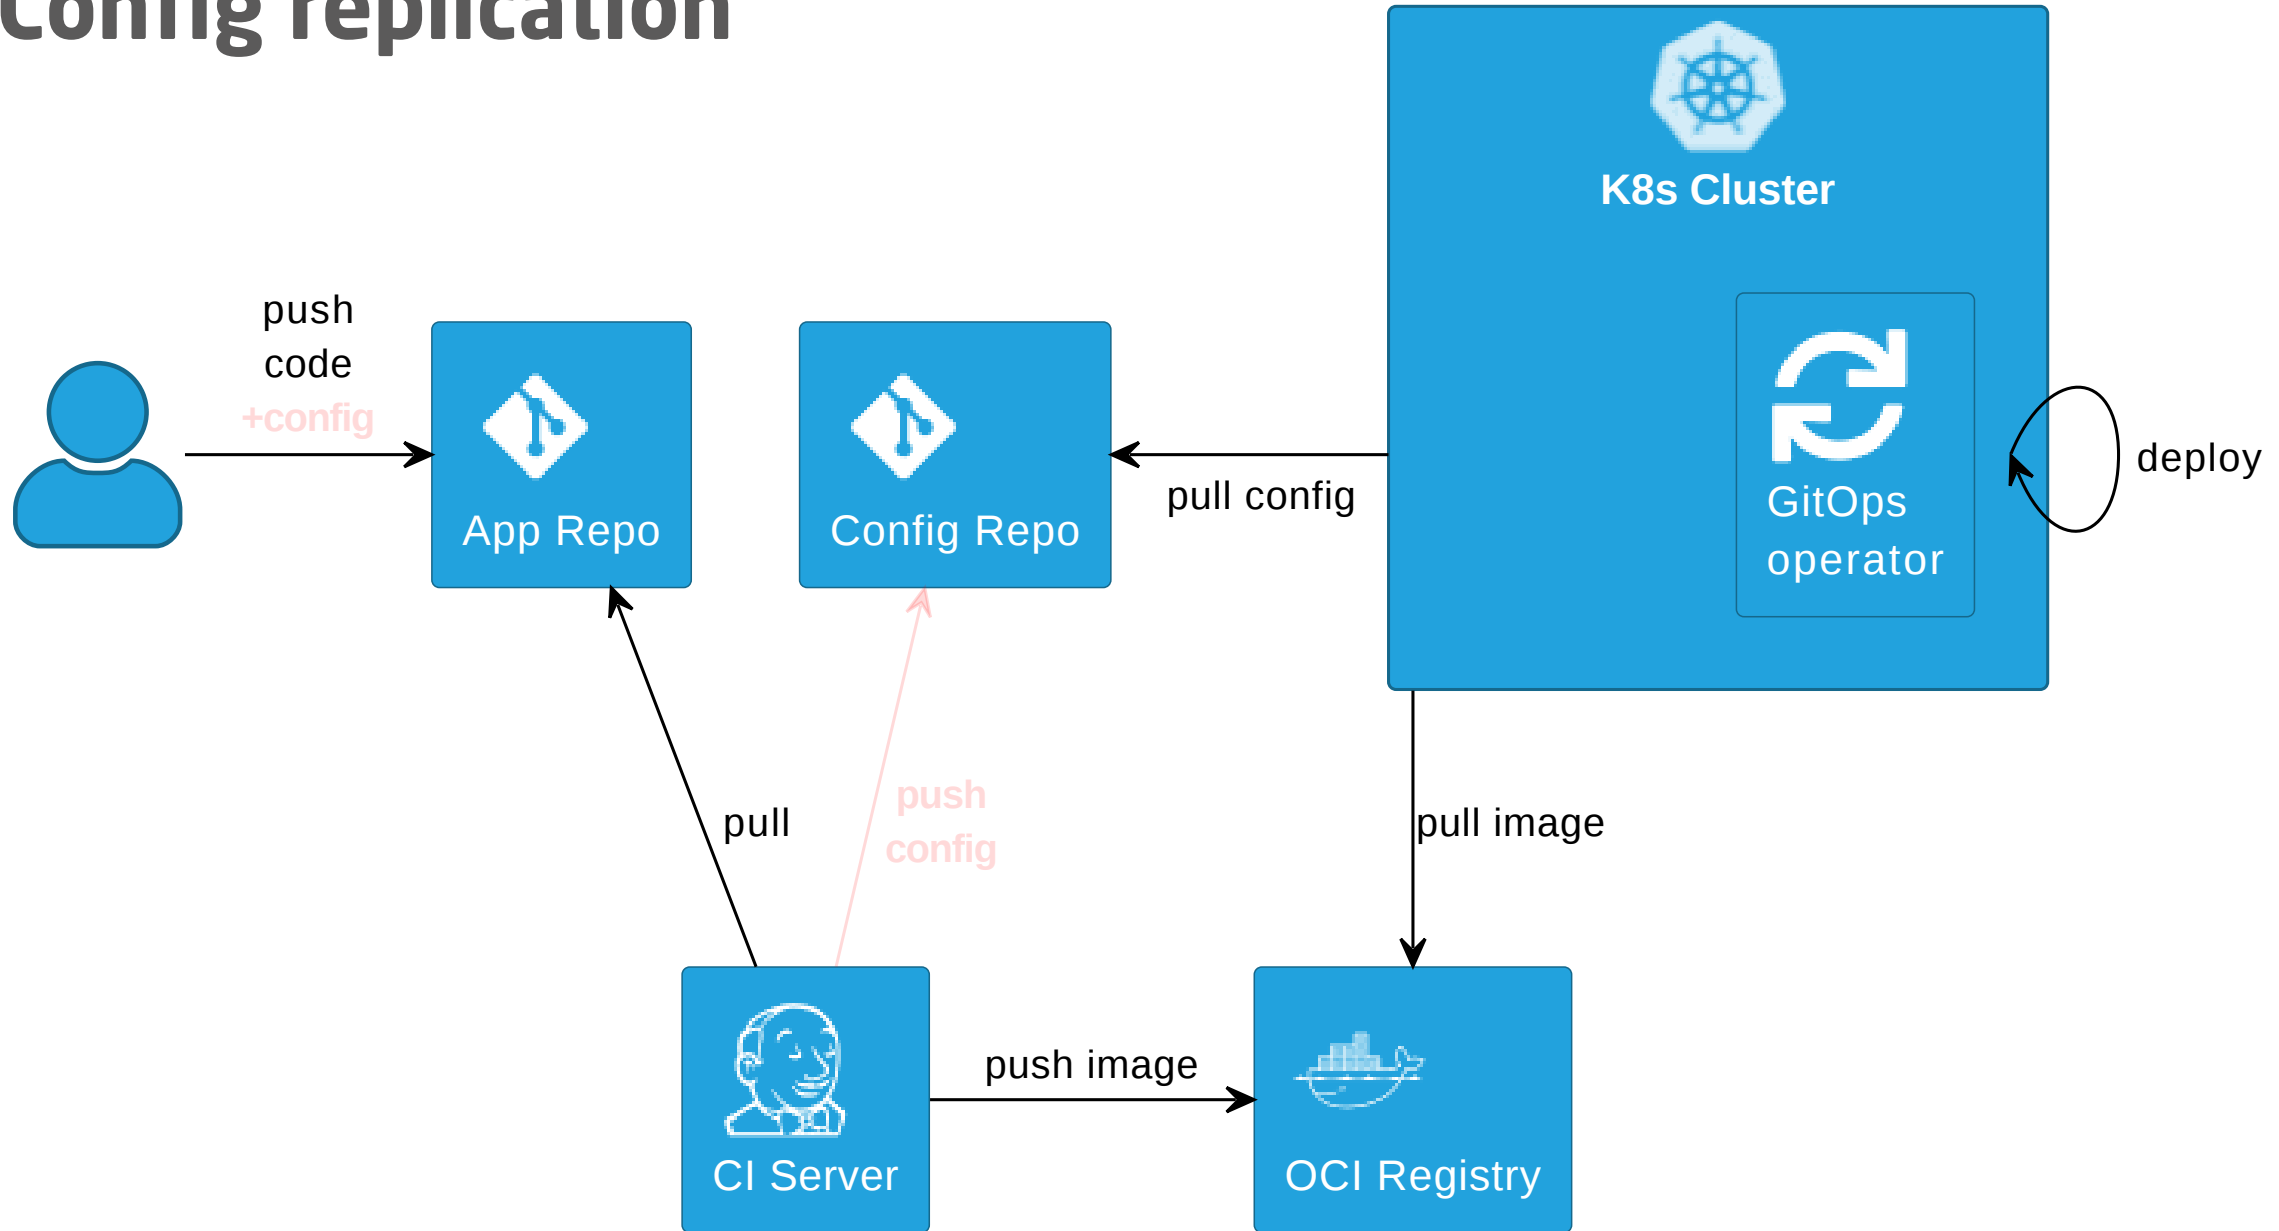
 argo-cd.readthedocs.io/en/release-2.8/user-guide/best_practices

Disadvantages

- Separated maintenance & versioning of app and infra code
- Review spans across multiple repos
- Local dev more difficult
- No static code analysis on config repo

How to avoid those?

Config replication



Advantages

- Single repo for development: higher efficiency
- Shift left: static code analysis + policy check on CI server, e.g. yamllint, kubeconform, helm lint, conftest, security scanners
- Automate config update (image tag + PR creation) 🕒
- Simplify review by adding info to PRs

Comments Commits Diff



[production] #2 [argocd/petclinic-plain@f448c2b](#)

Changeset [c9e3bf1](#) was committed 5 minutes ago

Authored by [Johannes Schnatterer](#) and committed by





Details



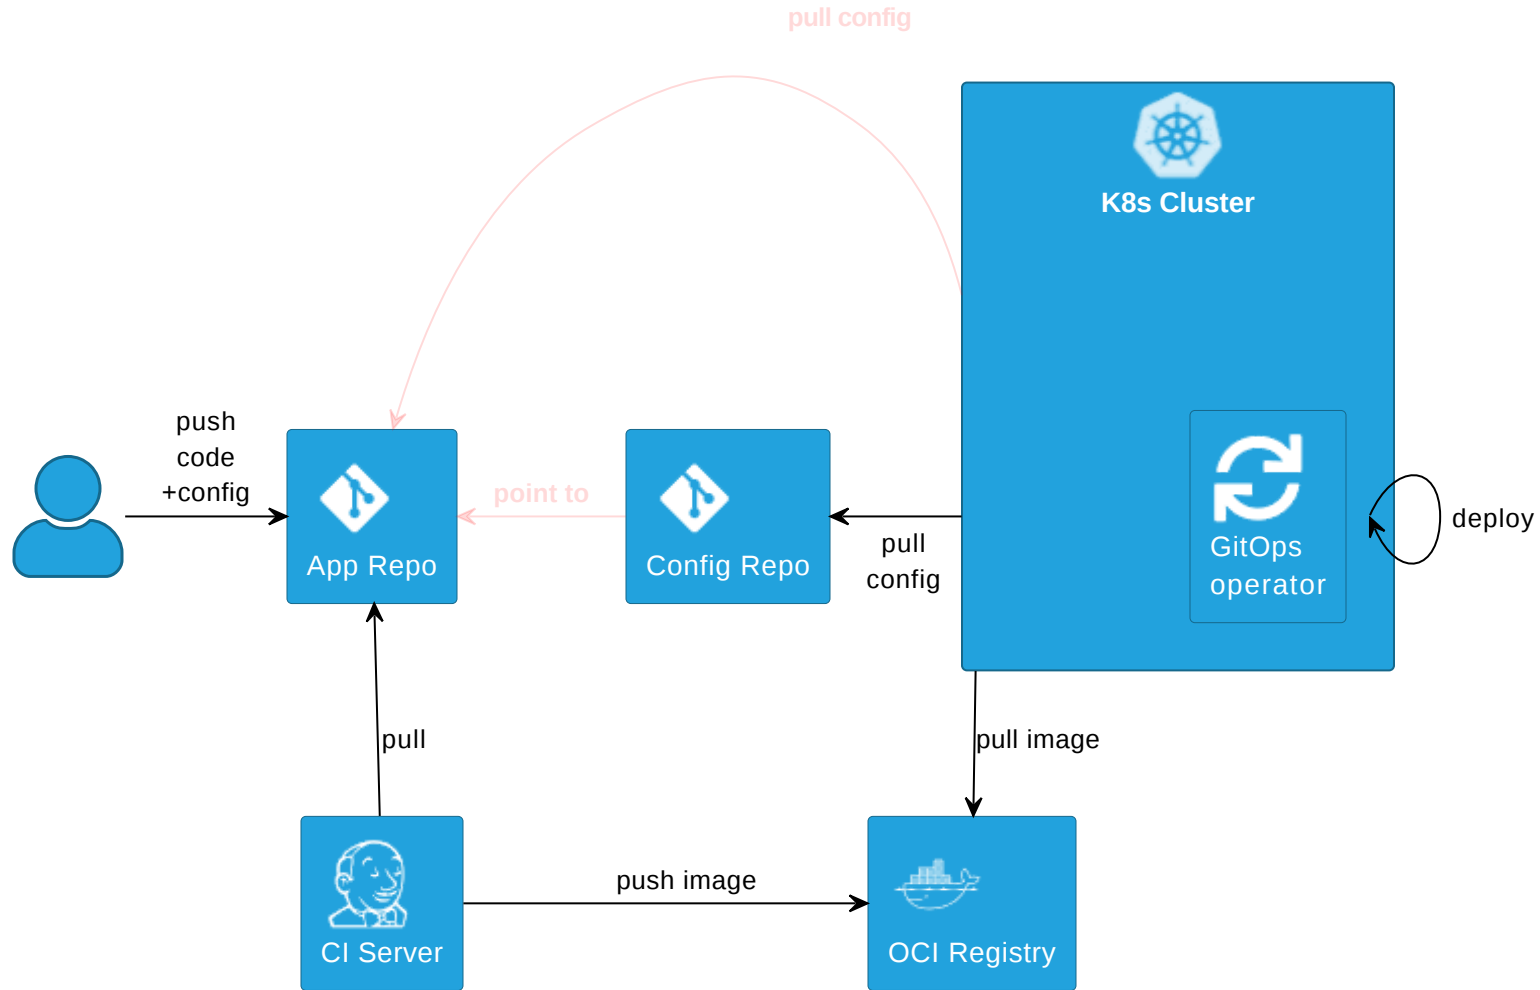
Sources



Disadvantages

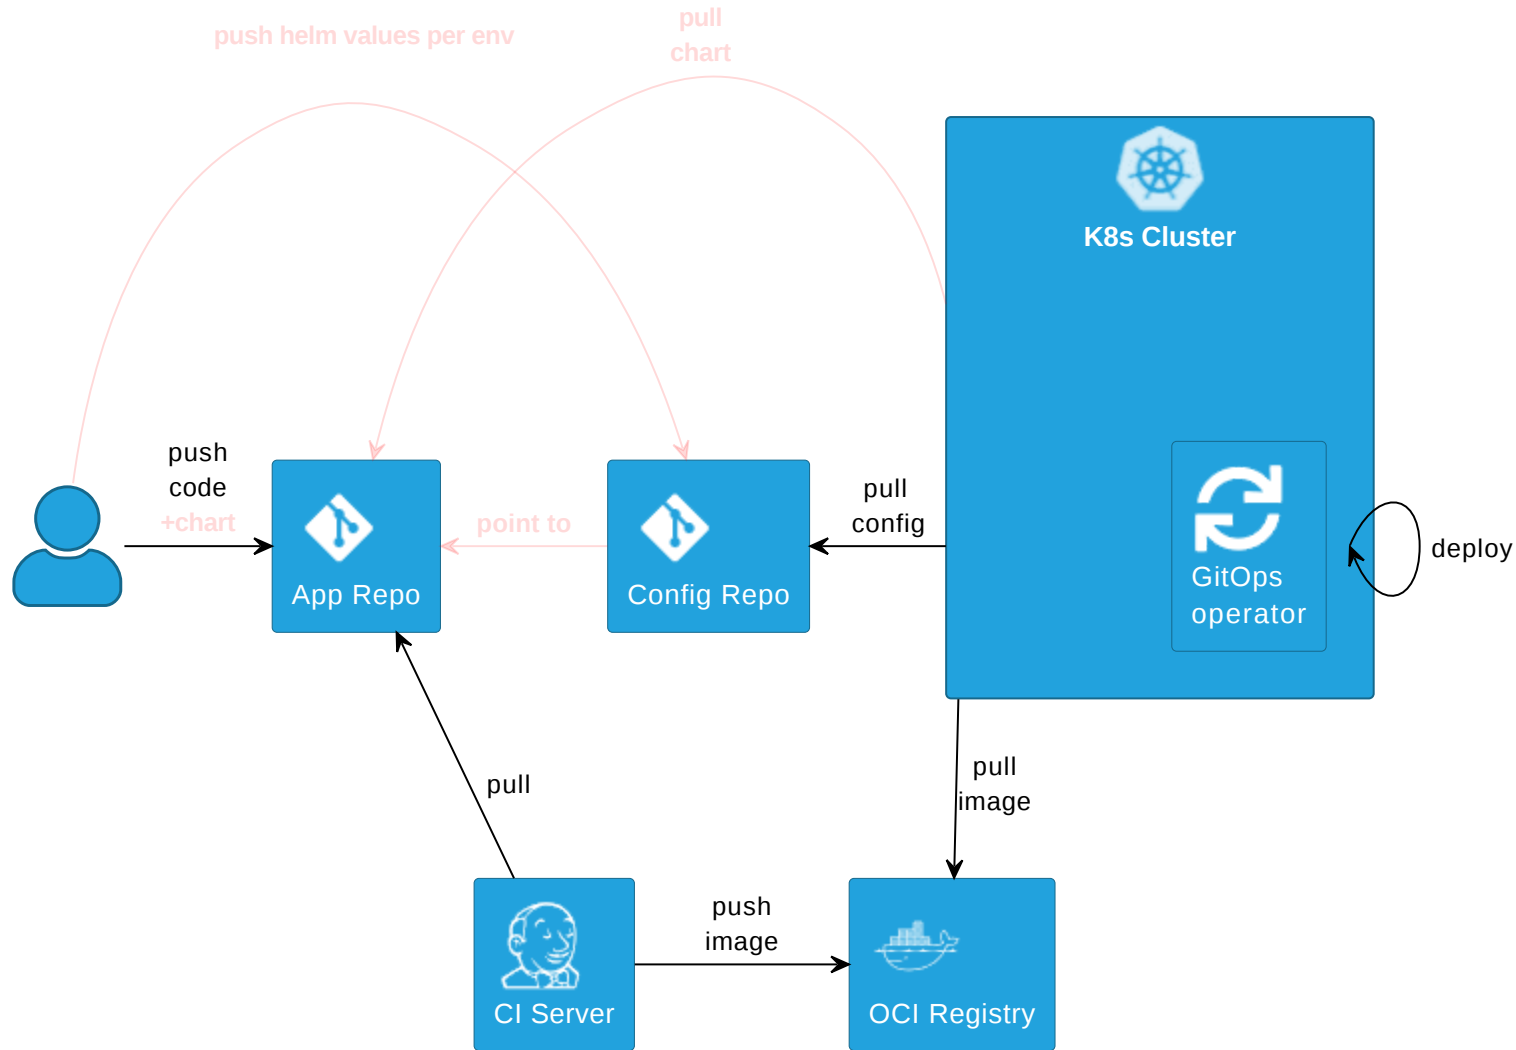
- Complexity in CI pipelines
 - ➔ Recommendation: Use a plugin or library, e.g.
 [cloudogu/gitops-build-lib](https://github.com/cloudogu/gitops-build-lib) 
- Redundant config (app repo + config repo)

Avoid Redundancy: Repo pointer



e.g.  fluxcd.io/flux/guides/repository-structure

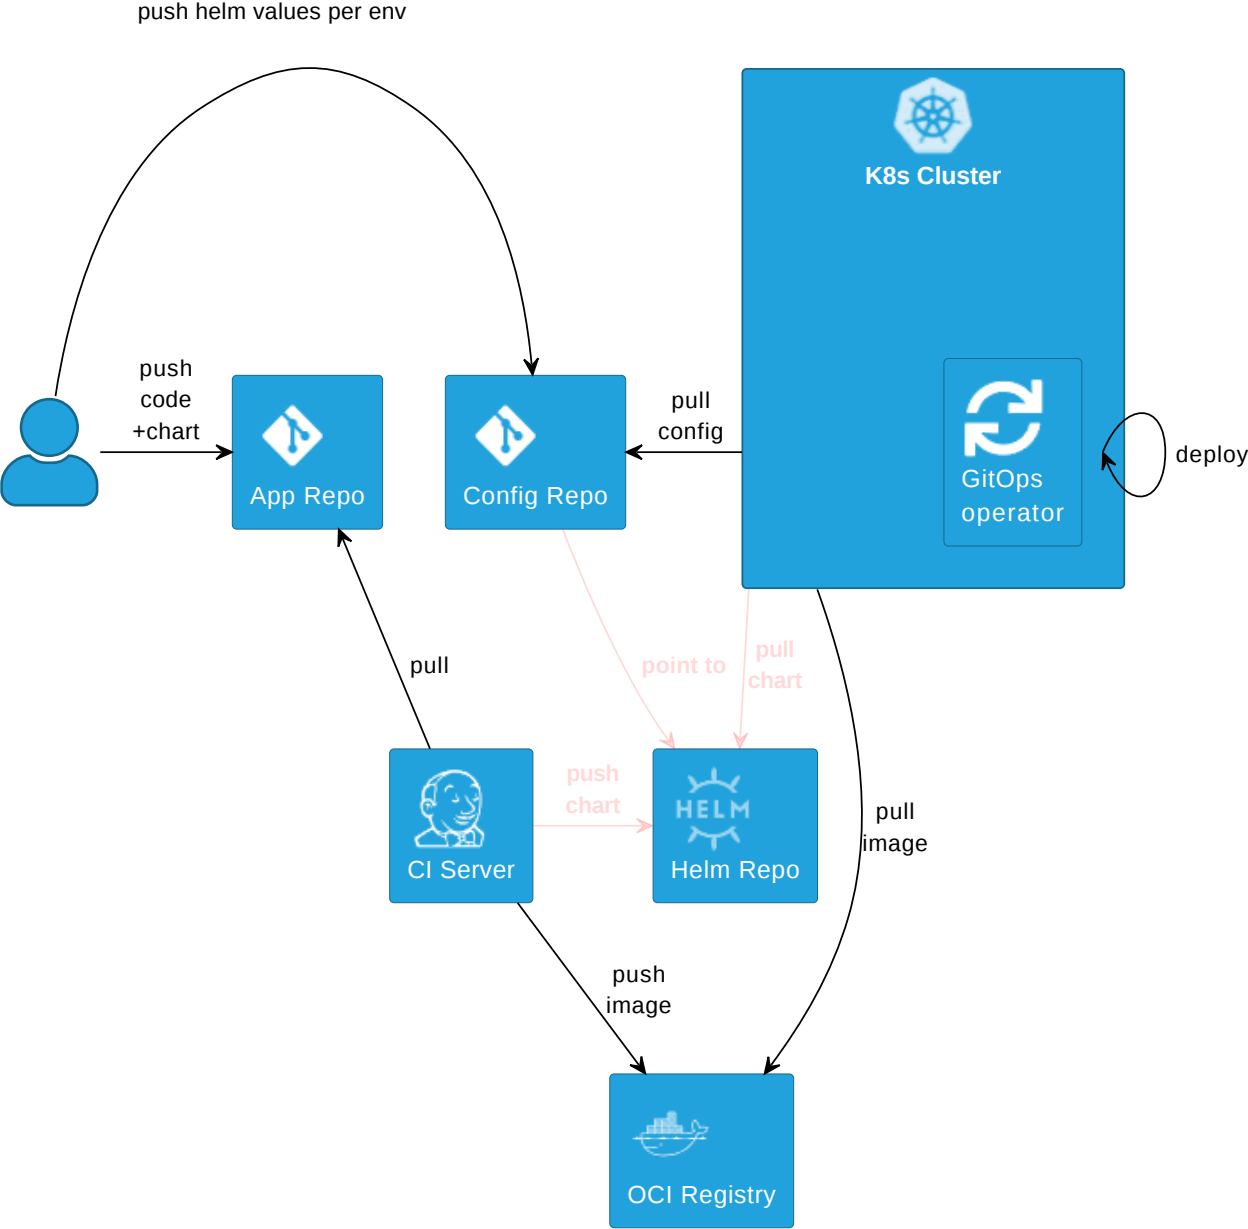
Middle ground: Config Split



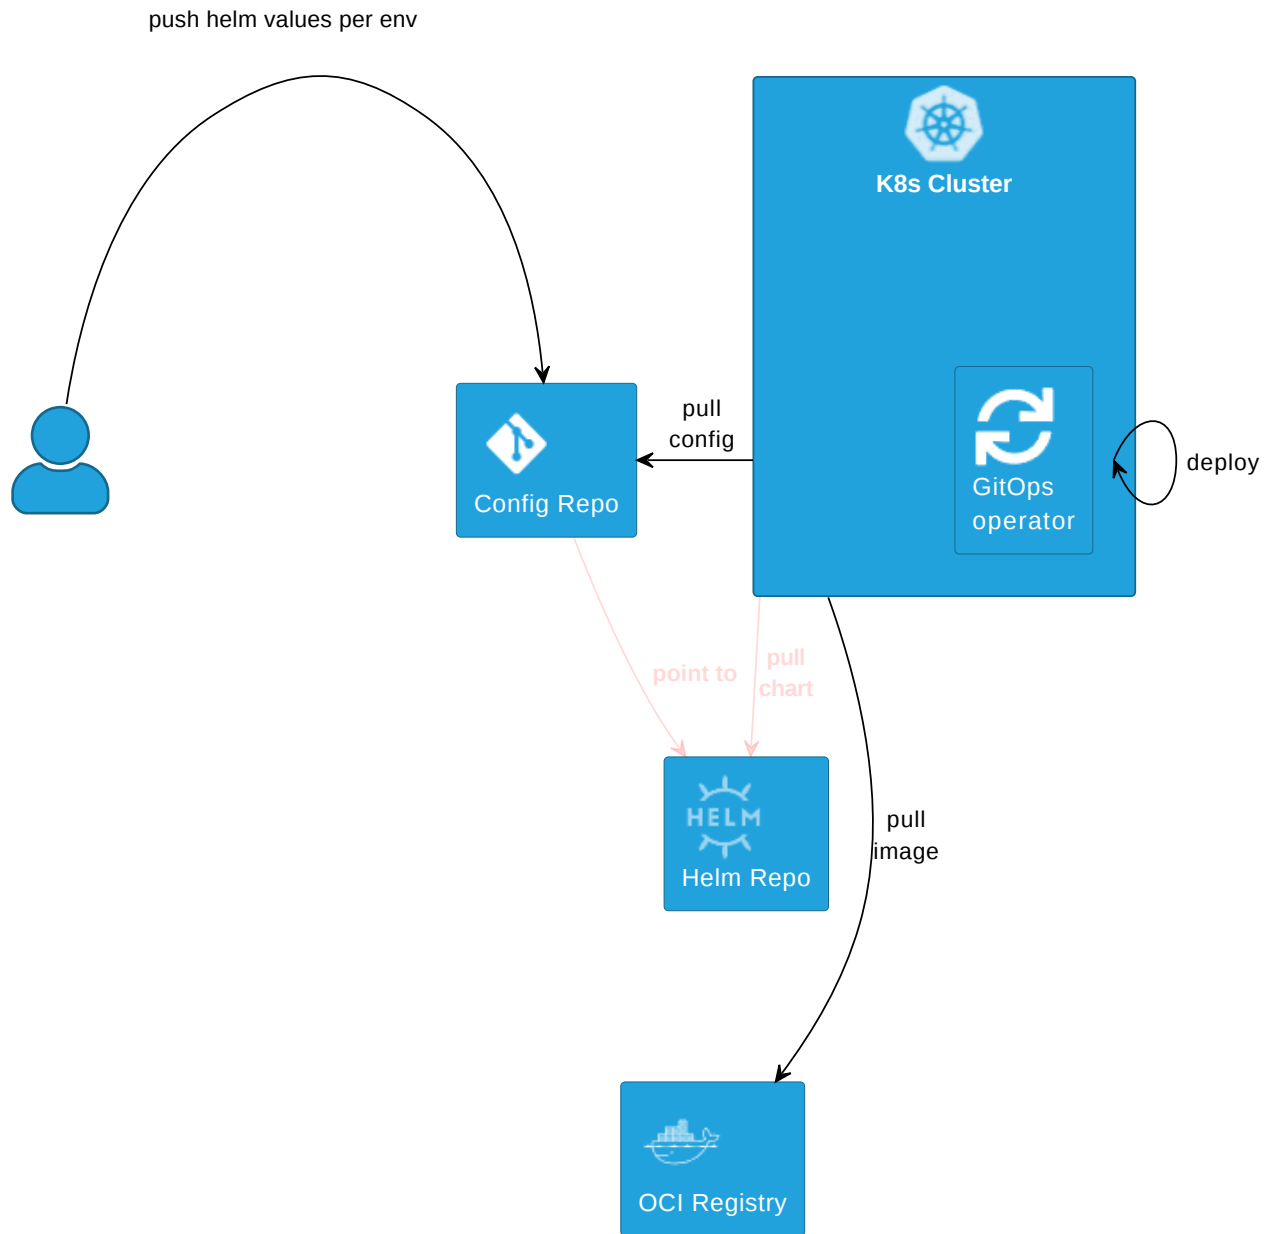
👉 HELM example

💡 Also works with K

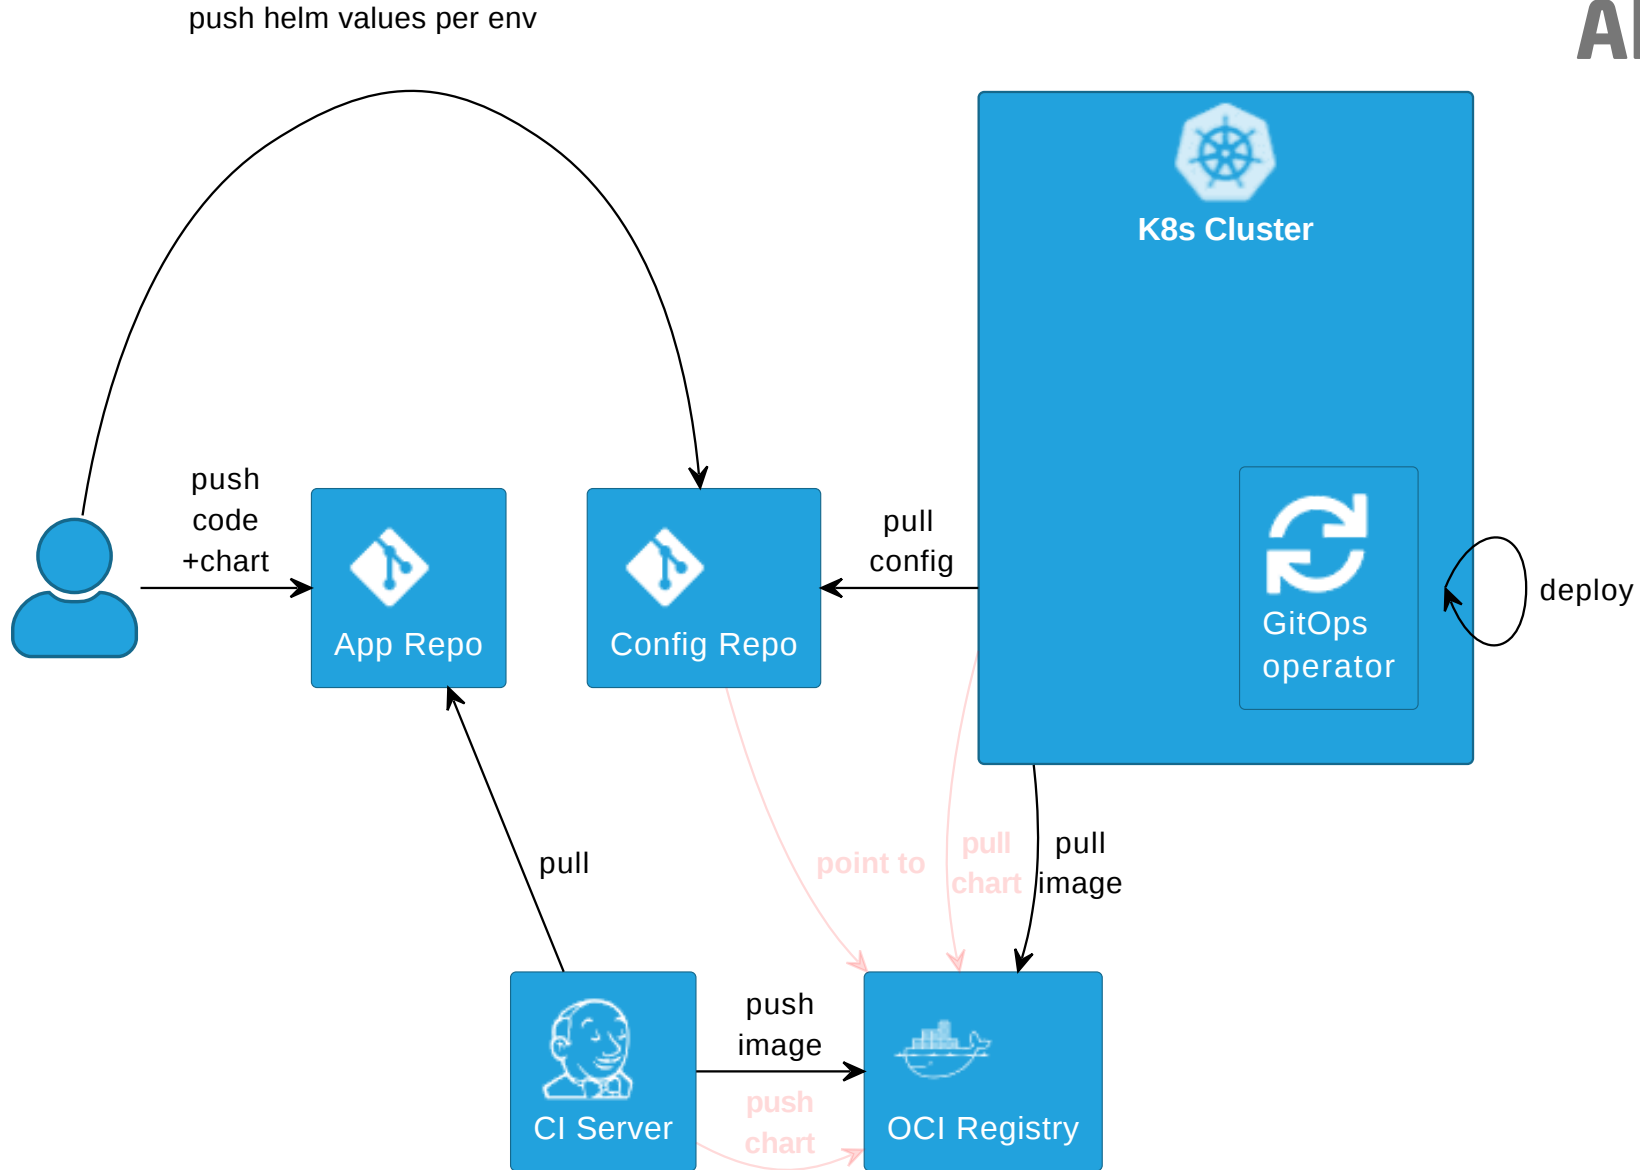
Alternative: Helm repo




💡 Same pattern for 3rd-party apps



Alternative 2: Helm in OCI




Alternative 3: OCI artifacts



 fluxcd.io/flux/cheatsheets/oci-artifacts

Promotion patterns









How to model environments/stages?

- **Branch per environment**
- **Folder/Directory per environment**
- **Repo per environment** (edge case)
-  **Preview environments**

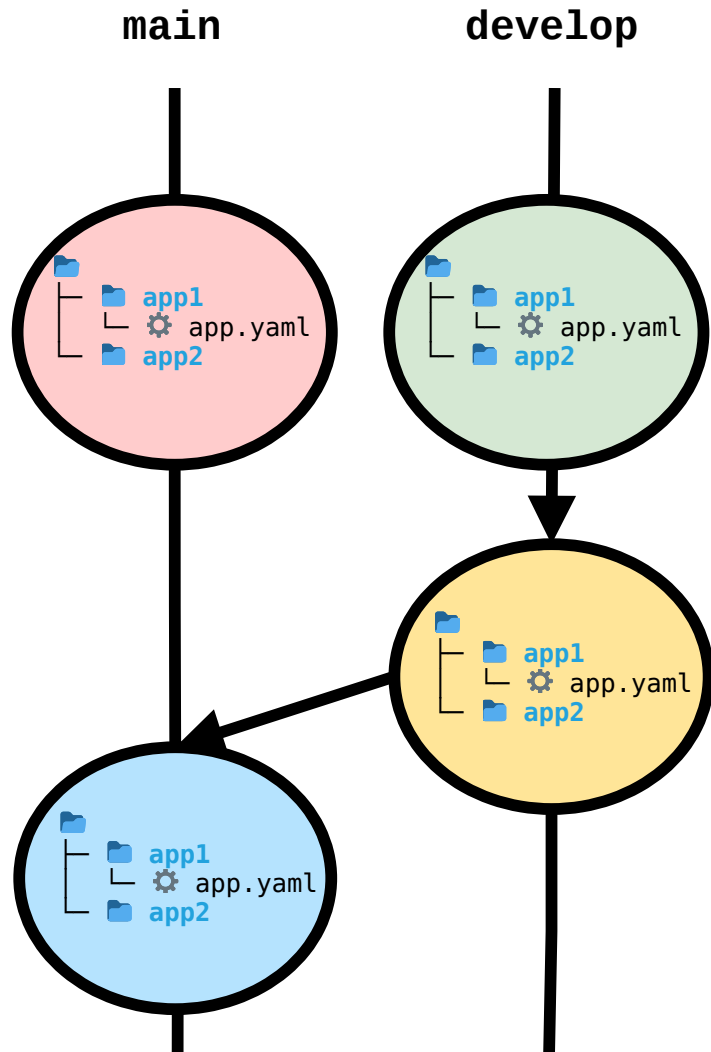
Branch vs folder per Environment

Branch per env

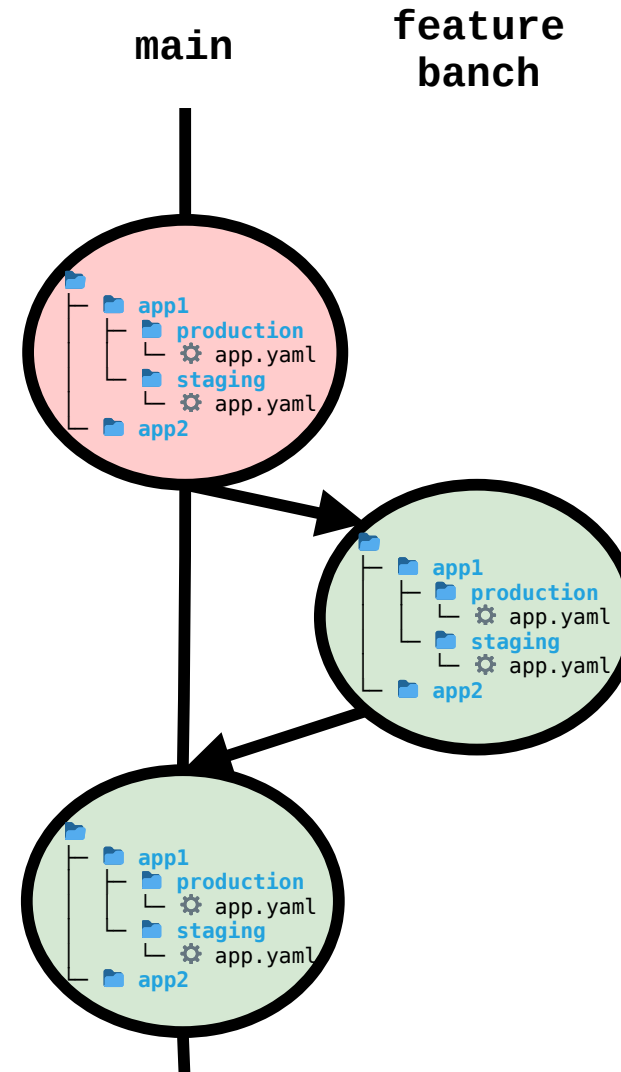
Folder per env

envs	permanent branches	trunk-based folders
mapping example	 develop  staging	 staging  Staging
	 main  production	 production  Production
promotion	merge	copy (+merge short-lived branches)

Branch per env



Folder per env



Branch per env

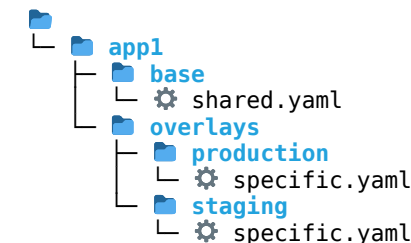
Folder per env

pros

- Forces PRs
- Feels natural for devs

- Avoids conflicts/drift
- Copy vs cherry pick
- Scales with envs

CM tool support (DRY)



references

1

1, 2, 3, 4, ...

Branches = anti-pattern

Repo per environment

Why would you want to use one repo per env?

- Access to folders more difficult to constrain than repos
- Organizational constraints, e.g.
 - "devs are not allowed to access prod"
 - security team needs to approve releases

 Repos more complicated than folders. Use only when really necessary.

Preview environments

- An environment that is created with a pull request
- and deleted on merge/close

 `ApplicationSet`, using the `PullRequest` generator

 `GitOpsSets` \approx 









Implementing promotion






Configuration Management tools

Tools for separating config of envs, keeping them DRY

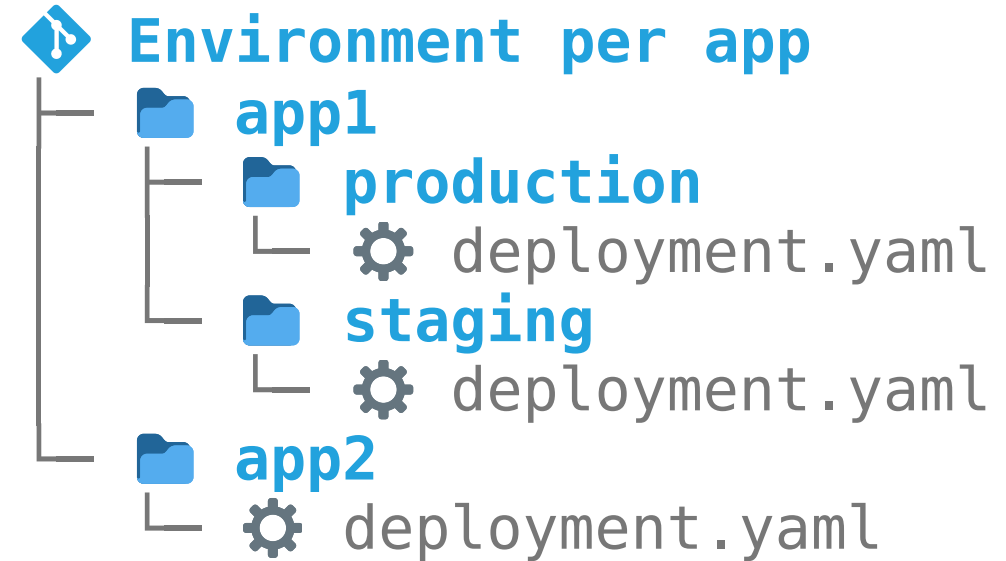
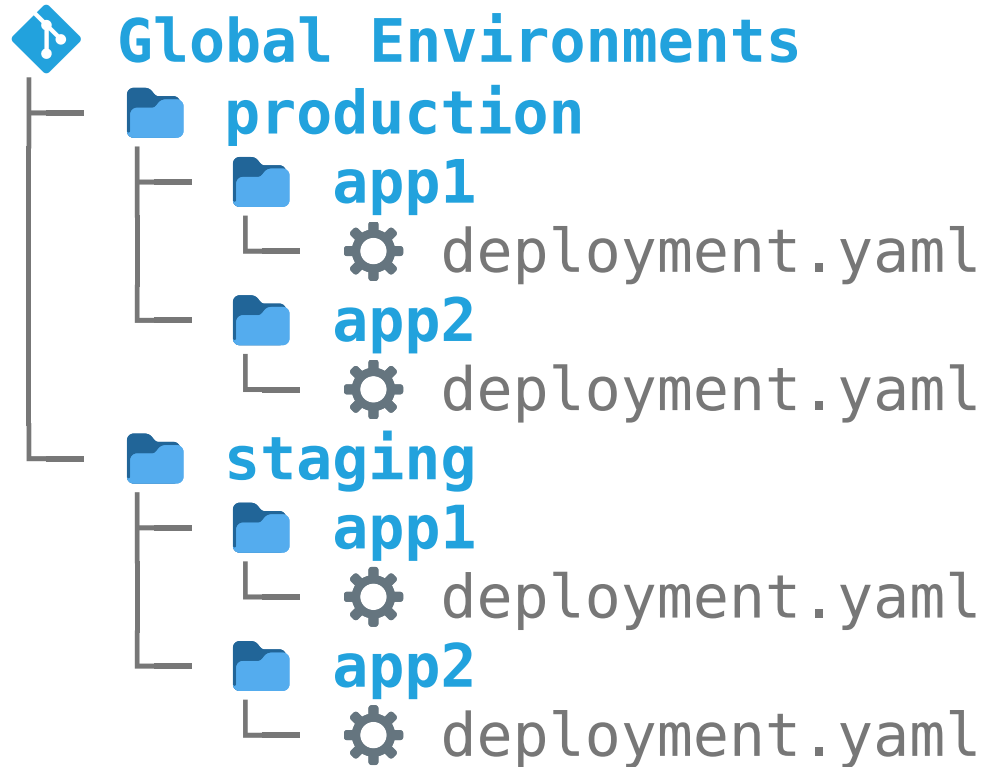
- Kustomize

- plain  `kustomization.yaml`    "agnostic"
- \neq Flux CRD  `Kustomization`
- `kustomize build / kubectl kustomize` via CI server 

- Helm

- CRD ( `Application`,  `HelmRelease`)
-  Umbrella Chart 
- `helm template` via CI server 

Global envs vs. env per app



e.g. **Preview Envs**

Config update

Who updates image version in config repo, creates branch and PR?

- **Manual:** Human pushes branch and create PR 🤡
- **CI Server:** Build job pushes branch, creates PR
- **Image Updater:** Operator pushes branch, create PR manually
- **Dependency Bot:** Bot pushes branch, creates PR

Promotion via CI Server

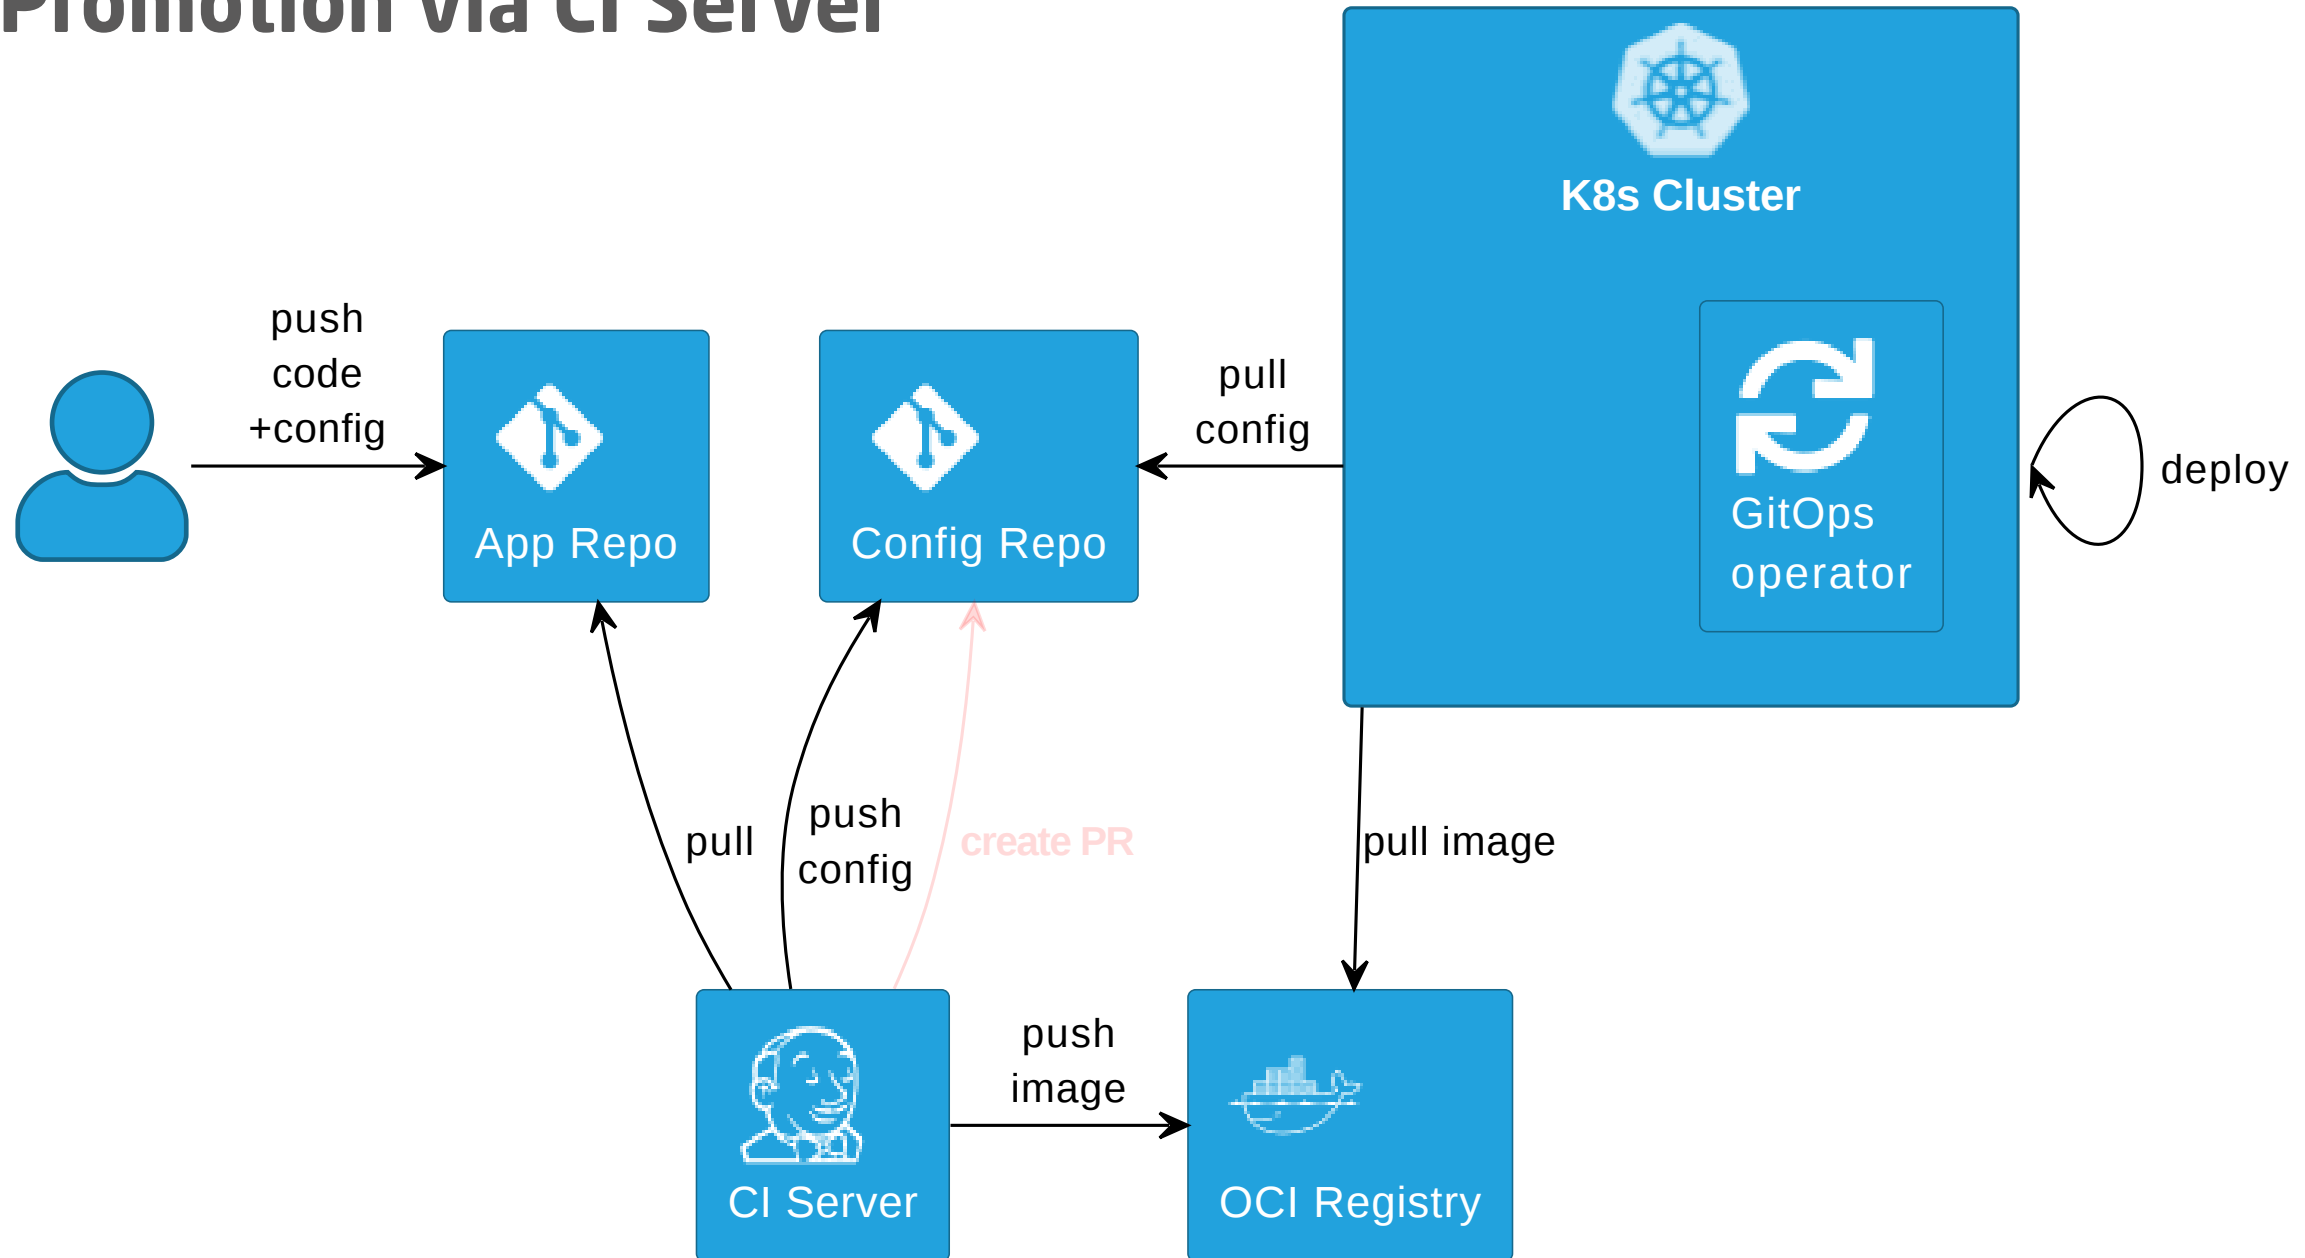
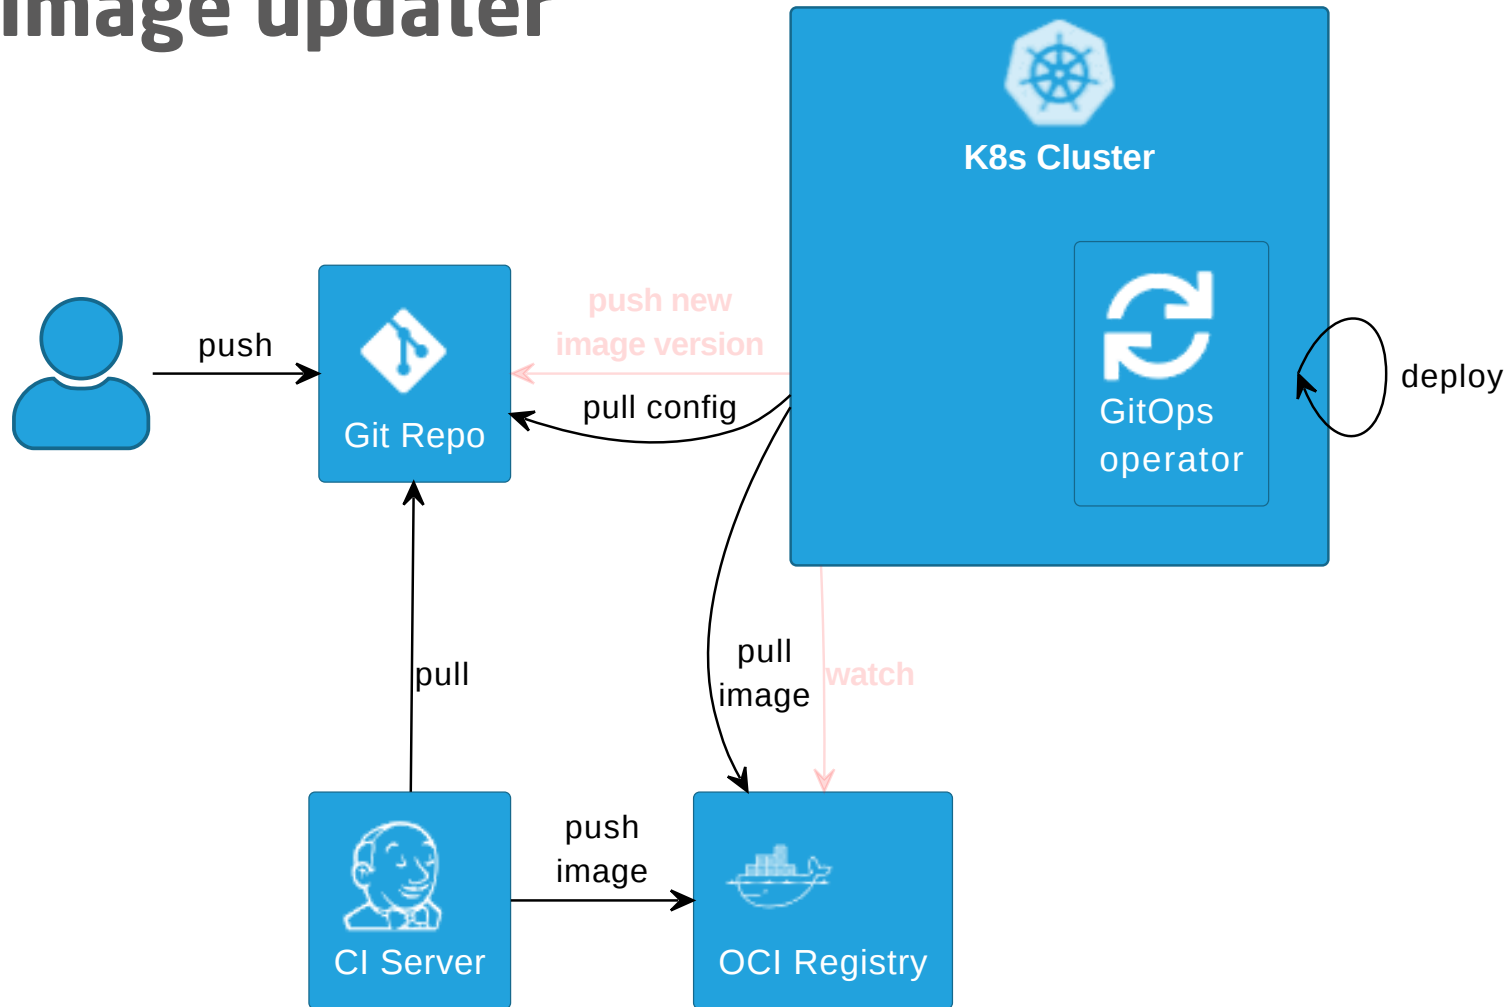


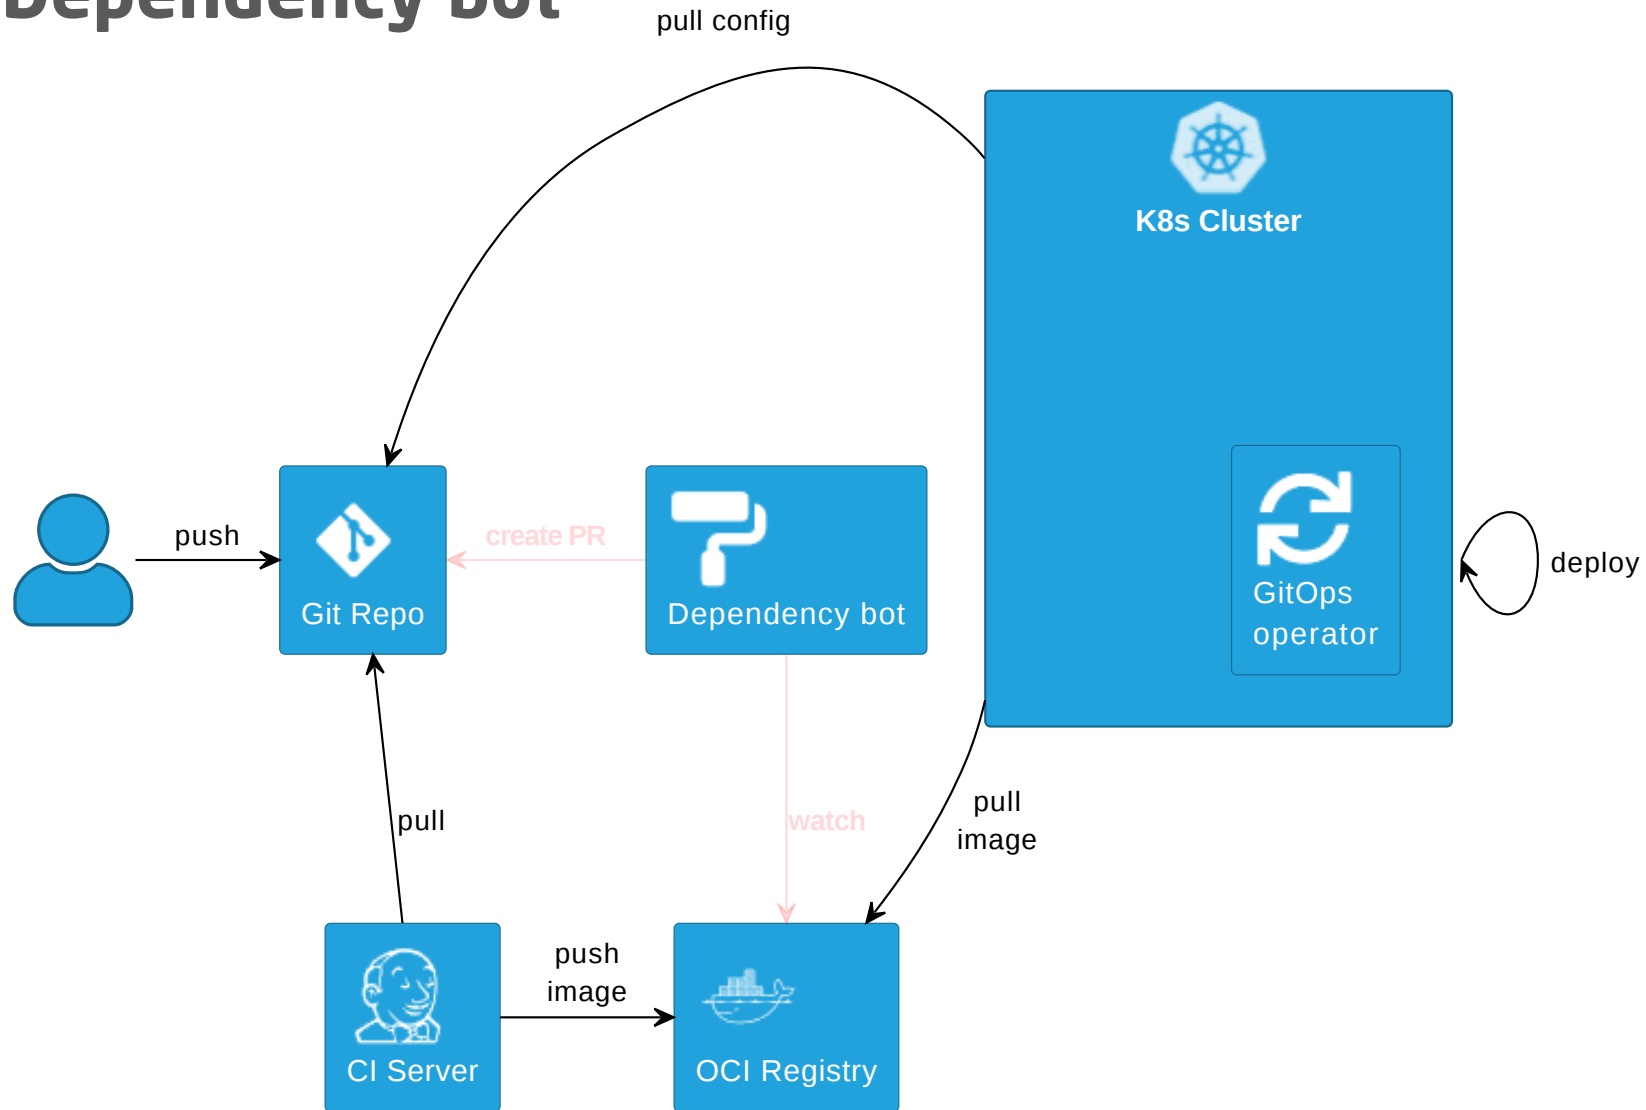
Image updater



 github.com/argoproj-labs/argocd-image-updater

 fluxcd.io/docs/guides/image-update

Dependency bot



e.g.  github.com/renovatebot/renovate

Pull Requests

GitOps - Operations by Pull Request

 weave.works/blog/gitops-operations-by-pull-request






But: avoid cargo cult

 PRs not mentioned in [principles](#)



Wiring patterns

Wiring up operator, repos, folders, envs, etc.

- **Bootstrapping:** `kubectl`, operator-specific CLI
- **Linking/Grouping:**
 - Operator-specific CRDs
 -  `Kustomization`
 -  `Application`
 - **Nesting:**  App of Apps
(same pattern with  `Kustomization`)
 - **Templating:**  `ApplicationSets` - folders, lists, config files, PRs

Example + Demo

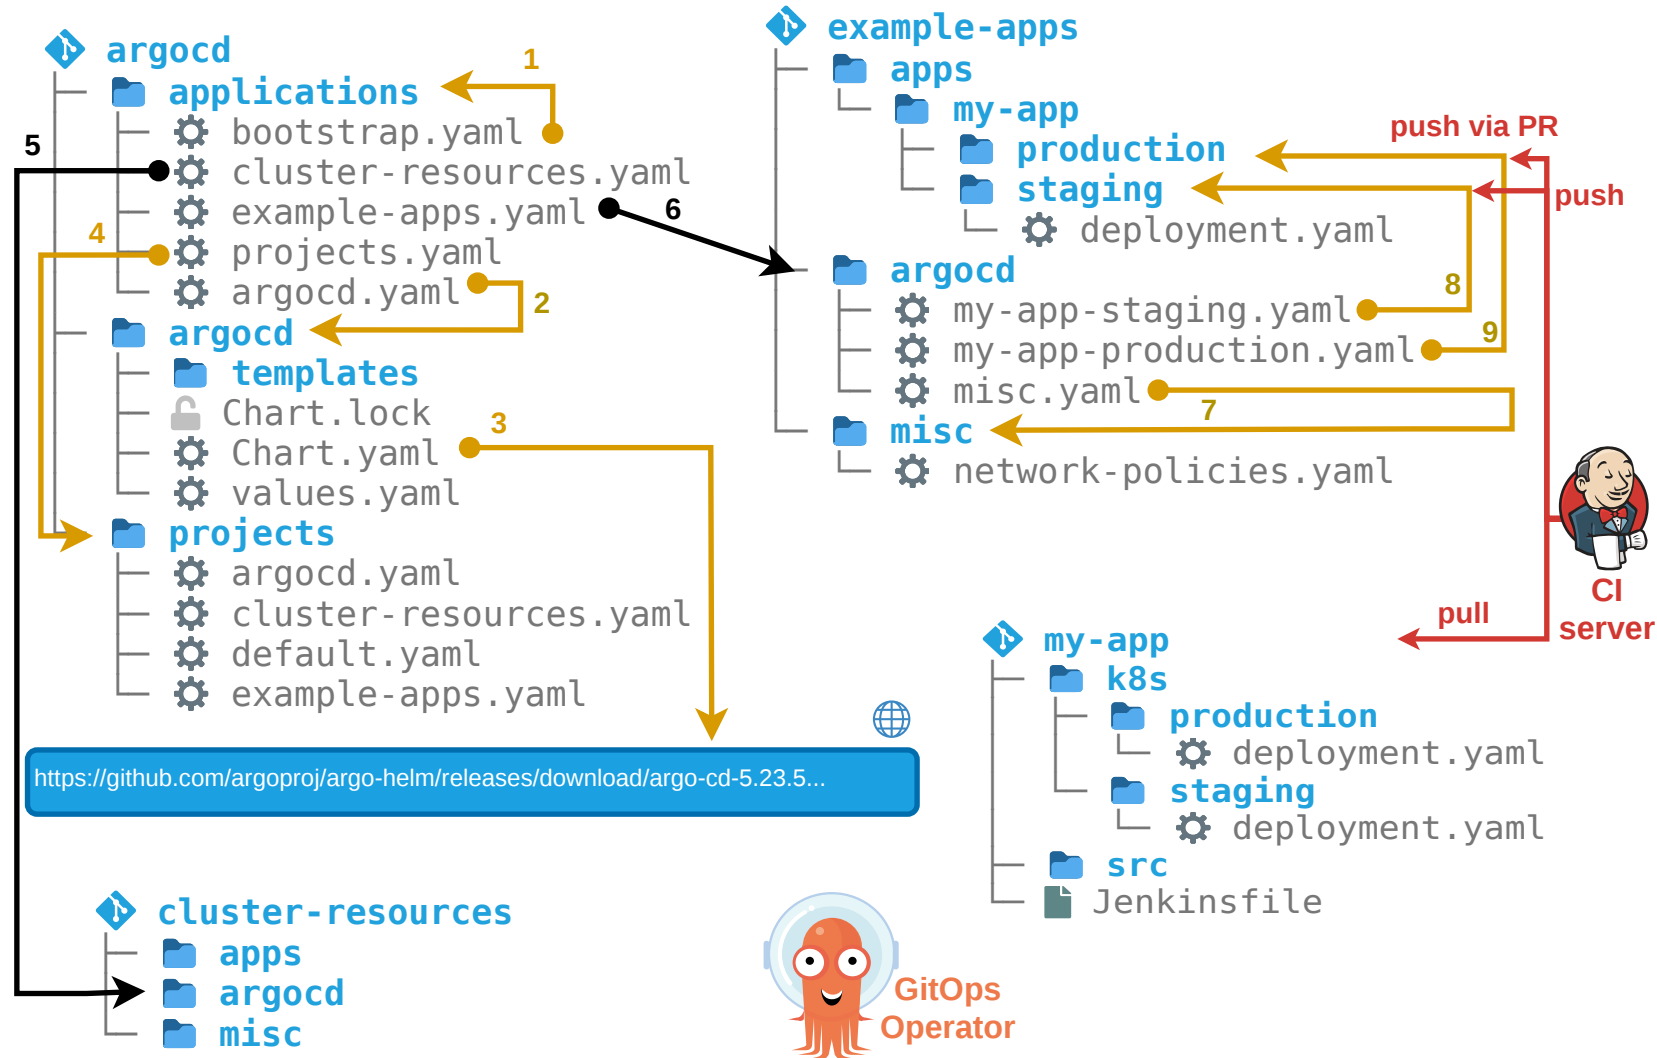
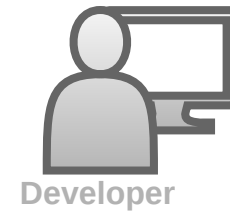


GitOps playground

- **Repo pattern:** Per team 🏢 per app
- **Operator pattern:** Instance per Cluster (Hub and Spoke)
- **Operator:** 🧙
- **Boostrapping:** Helm, kubectl
- **Linking:** 🧙 Application
- **Features:**
 - Operate ArgoCD with GitOps
 - Solution for cluster resources
 - Config update + replication via CI
 - Mixed repo patterns
 - Env per app pattern
- **Source:** 🔄 [cloudogu/gitops-playground](https://github.com/cloudogu/gitops-playground)

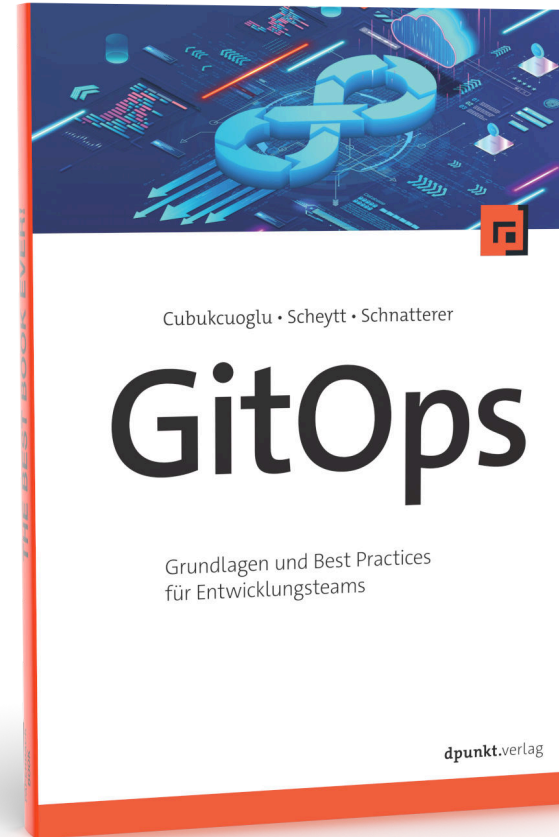
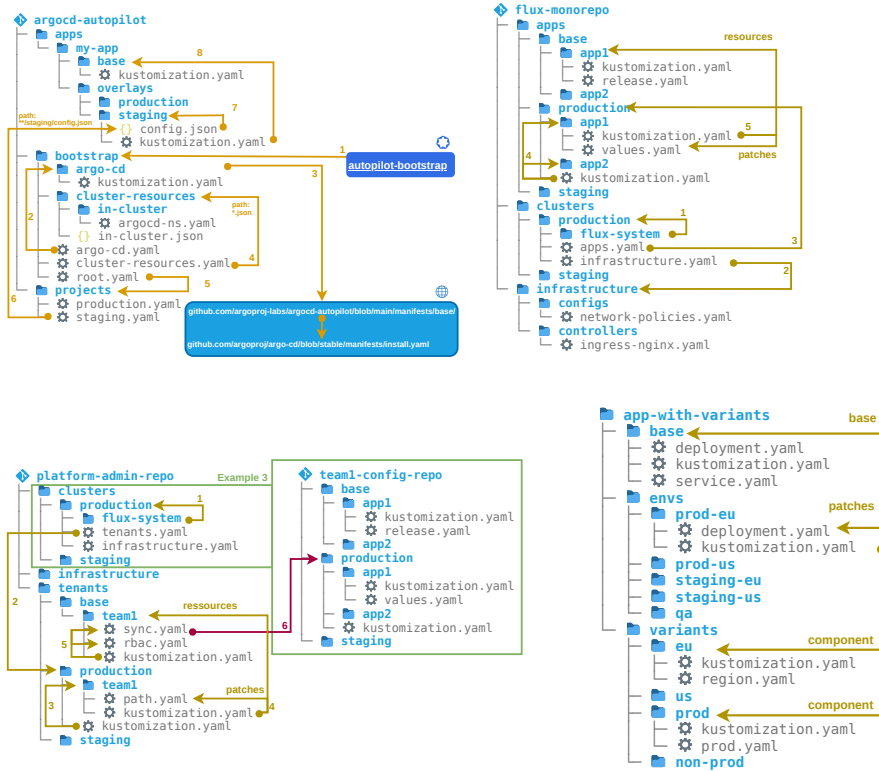


```
COMMIT='8e21bd4'
bash <(curl -s \
  "https://raw.githubusercontent.com/cloudogu/gitops-playground/$COMMIT/scripts/init-cluster.sh") \
  --bind-ingress-port=80 \
  && sleep 2 && docker run --rm -it --pull=always -u $(id -u) \
    -v ~/.config/k3d/kubeconfig-gitops-playground.yaml:/home/.kube/config \
    --net=host \
    ghcr.io/cloudogu/gitops-playground:$COMMIT --yes --argocd --base-url=http://local.gd
# Other interesting options --monitoring --vault=dev
```

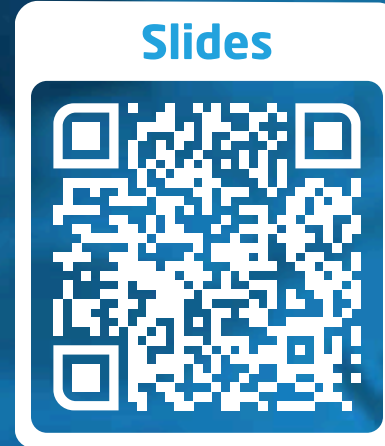


More examples + further reading

 [cloudogu/gitops-patterns](https://github.com/cloudogu/gitops-patterns)



Johannes Schnatterer, Cludogu GmbH



💪 Join my team: cludogu.com/join/cloud-engineer

📧 @schnatterer@floss.social

🌐 in/jschnatterer

🐦 @jschnatterer

Wir entwickeln einen open source GitOps-Stack für K8s

Sag uns wie wir GitOps für dich leichter machen können



Image sources

- implementation
<https://unsplash.com/photos/selective-focus-photography-blue-and-black-makita-power-drill-KlbyOnxseY8>
- Demo <https://unsplash.com/photos/assorted-color-hot-air-balloons-during-daytime-DuBNA1QMpPA>
- coloured-parchment-paper background by brgfx on Freepik
https://www.freepik.com/free-vector/coloured-parchment-paper-designs_1078492.htm
- Question mark
<https://pixabay.com/illustrations/question-mark-question-response-1020165/>